Knowledge Discovery in Databases (KDD) Methodology

Applied if the system involves data analytics or forecasting. It follows stages from raw data to usable patterns.

What data?

The primary dataset for this study consists of historical transaction records obtained from the POS system of Banelo Bake and Brew. The data includes essential fields such as product names, categories, quantities sold, unit prices, total sales per transaction, inventory levels, and transaction timestamps. These attributes were chosen for their direct relevance to sales behavior and forecasting accuracy. Although promotional data is not available, time-based attributes like weekdays, weekends, and months are considered to capture periodic fluctuations in customer demand.

Dataset Characteristics:

Total Records: 3,311 sales transactions
Time Span: 407 days of historical data
Date Range: Comprehensive coverage of daily sales patterns

Data Attributes:

Table #

Data Attributes for Sales Forecasting Model

Attribute	**Data Type**	**Description**
Product Name	String	The specific item sold (e.g., "Iced Coffee", "Blueberry Muffin")
Category	Categorical	Product classification (Beverage or Pastry)
Quantity Sold	Integer	Number of units sold per transaction
Price	Decimal	Unit price at the time of sale (PHP)

Date	DateTime	Transaction timestamp
Revenue	Decimal	Total sales amount (Quantity × Price)

Note. This table presents the input features used in the machine learning model for sales prediction.

How did I prepare it?

Data preparation began with preprocessing, which involves cleaning the raw dataset to ensure consistency, completeness, and accuracy. This includes removing duplicate entries, correcting inconsistent product codes or labels, standardizing date formats, and addressing any missing or null values in key fields. After cleaning, the data underwent a transformation to make it suitable for mining. This involves aggregating daily sales into weekly or monthly totals to smooth short-term noise, generating lag features (e.g., sales from previous days or weeks), and creating rolling averages to highlight trends. Additionally, categorical fields like product categories were encoded numerically using techniques such as one-hot encoding. Time-based features such as day of the week, month, or whether the date falls on a weekend, is also engineered to help the model detect seasonal and behavioral sales patterns.

What model I used?

To ensure the selection of the most accurate forecasting model, a systematic comparative evaluation was conducted across six (6) different machine learning algorithms. This rigorous approach aligns with best practices in predictive analytics and provides empirical justification for model selection.

Table #

Machine Learning Algorithms Evaluated

#	Algorithm	Type	Description
1	Linear Regression	Statistical Model	Baseline model using least squares

			regression
2	Decision Tree	Tree-Based Model	Non-linear model using recursive partitioning
3	Random Forest	Ensemble (Bagging)	Ensemble of decision trees using bootstrap aggregating
4	Gradient Boosting	Ensemble (Boosting)	Sequential ensemble correcting previous errors
5	XGBoost	Optimized Boosting	Optimized gradient boosting with regularization
6	LightGBM	Fast Boosting	Gradient boosting optimized for efficiency

Note. Six machine learning algorithms were evaluated to identify the optimal forecasting model for the Banelo system.

Model Training Methodology

All models were trained using a consistent, fair comparison methodology:

Training Set: 80% of historical data (oldest 2,649 records)
Testing Set: 20% of historical data (most recent 662 records)
Validation Strategy: Temporal split to prevent data leakage and simulate real-world forecasting
Feature Set: Identical 10+ engineered features used across all models
Random Seed: Fixed seed (42) for reproducibility

Evaluation Metrics

Model performance was assessed using four industry-standard regression metrics:

1. Mean Absolute Error (MAE)

- Measures the average absolute difference between predicted and actual values
- Expressed in the same units as the target variable (quantity sold)
- Lower values indicate better accuracy

MAE = $(1/n) \times \Sigma|y_i - \hat{y}_i|$

Where:

- n = number of observations
- y_i = actual value
- $\hat{y}_i$ = predicted value
- $|...|$ = absolute value

2. Root Mean Squared Error (RMSE)

- Penalizes larger errors more heavily than MAE due to squaring
- Provides insight into model robustness against large prediction errors
- More sensitive to outliers than MAE

RMSE = $\sqrt{[(1/n) \times \Sigma(y_i - \hat{y}_i)^2]}$

Where:

- n = number of observations
- y_i = actual value
- $\hat{y}_i$ = predicted value
- $\sqrt{}$ = square root

3. Mean Absolute Percentage Error (MAPE)

- Expresses prediction error as a percentage of actual values
- Scale-independent metric enabling comparison across different datasets

- Lower percentages indicate higher forecasting accuracy

MAPE = $(100/n) \times \Sigma|(y_i - \hat{y}_i) / y_i|$

Where:

- n = number of observations
- y_i = actual value
- $\hat{y}_i$ = predicted value
- $|...|$ = absolute value
- Result is expressed as a percentage (%)

4. R² Score (Coefficient of Determination)

- Measures the proportion of variance in sales explained by the model
- Ranges from $-\infty$ to 1.0, with values closer to 1.0 indicating better explanatory power
- Negative values indicate predictions worse than the mean baseline

R^2 = $1 - (SS\_res / SS\_tot)$

Where:

- $SS\_res = \Sigma(y_i - \hat{y}_i)^2$ (Residual Sum of Squares)
- $SS\_tot = \Sigma(y_i - \bar{y})^2$ (Total Sum of Squares)
- y_i = actual value
- $\hat{y}_i$ = predicted value
- $\bar{y}$ = mean of actual values

Alternative formula:

$R^2 = 1 - [\Sigma(y_i - \hat{y}_i)^2 / \Sigma(y_i - \bar{y})^2]$

Experimental Results

Table #

Experimental Results: Comprehensive Model Comparison With Performance Metrics

Model	**MAE**	**RMSE**	**MAPE**	**R^2**
Linear Regression ✓	1.0623	1.2694	55.13%	0.1897
Random Forest	1.1123	1.3261	57.50%	0.1157
XGBoost	1.1189	1.3342	57.58%	0.1050
LightGBM	1.1191	1.3418	57.48%	0.0947
Gradient Boosting	1.1200	1.3426	57.45%	0.0937
Decision Tree ✗	1.2826	1.5823	64.81%	-0.2590

Note. ✓ = Best performing model; ✗ = Worst performing model. Lower values indicate better performance for MAE, RMSE, and MAPE, while higher values are better for R^2. MAE = Mean Absolute Error; RMSE = Root Mean Squared Error; MAPE = Mean Absolute Percentage Error; R^2 = Coefficient of Determination.

Figure #

Mean Absolute Error Comparison Across Machine Learning Models

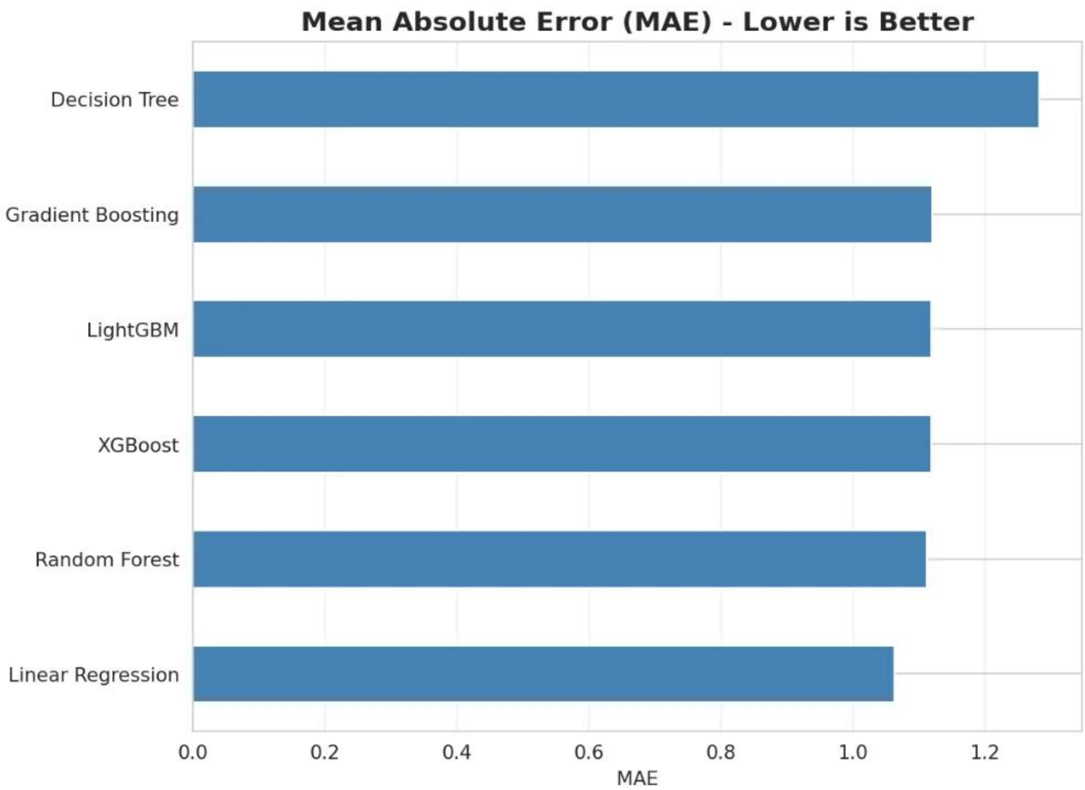

Note. This figure demonstrates the MAE values for six algorithms. Linear Regression achieved the lowest error at 1.06 units, indicating superior predictive accuracy. The x-axis represents the MAE values in units, while the y-axis lists the machine learning models evaluated. Lower MAE values indicate better performance in minimizing average prediction errors.

Figure #

Root Mean Squared Error Comparison Across Machine Learning Models

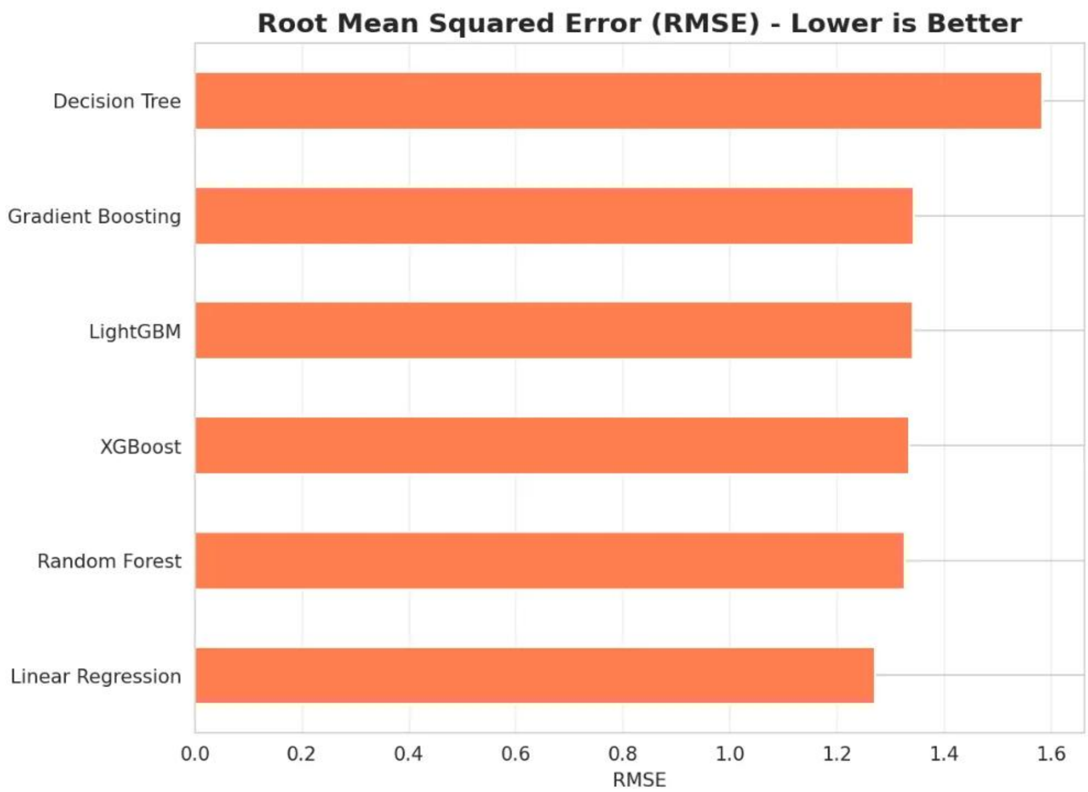

Note. This figure compares RMSE values across six algorithms, with Linear Regression recording 1.27 units. RMSE penalizes larger deviations more heavily than MAE, confirming Linear Regression's consistent accuracy. The x-axis shows RMSE values in units, and the y-axis displays the models. The lower RMSE of Linear Regression indicates minimal large prediction errors.

Figure #

Mean Absolute Percentage Error Comparison Across Machine Learning Models

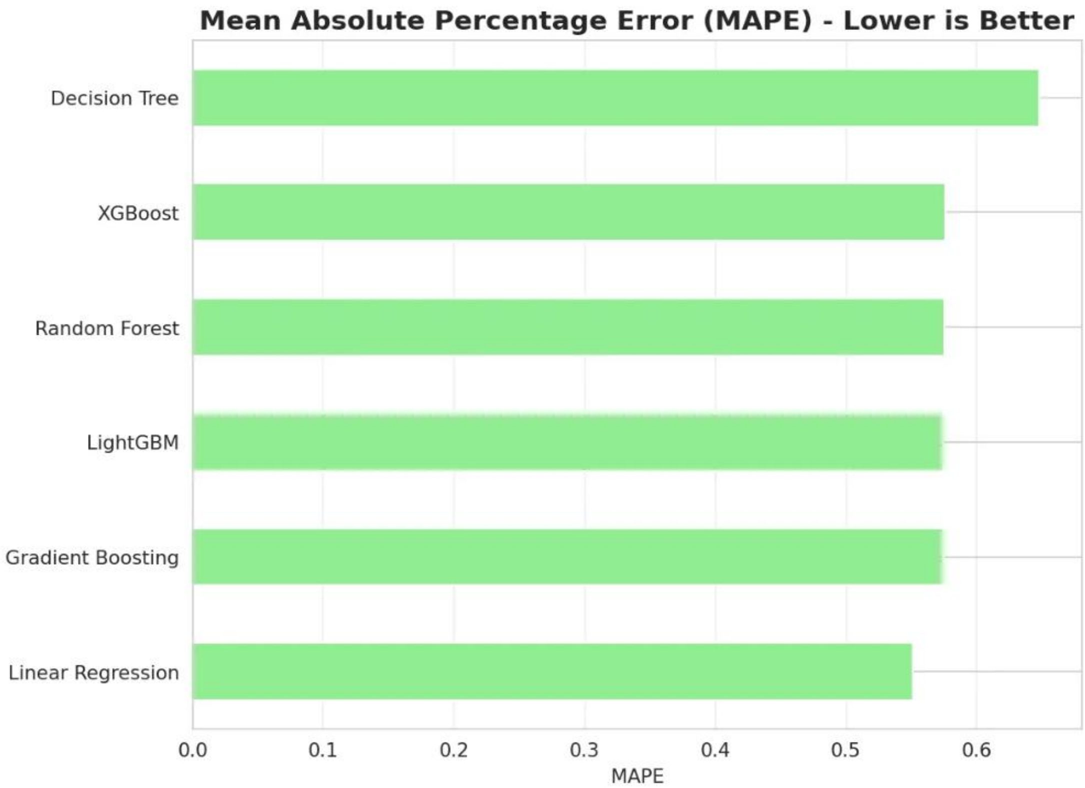

Note. This figure displays MAPE values for all evaluated models, with Linear Regression achieving 55.13%. The x-axis represents MAPE as a percentage, while the y-axis lists the algorithms. While MAPE indicates moderate percentage deviation, the absolute error metrics (MAE, RMSE) demonstrate practical forecasting accuracy for inventory management decisions.

Figure #

Coefficient of Determination (R²) Comparison Across Machine Learning Models

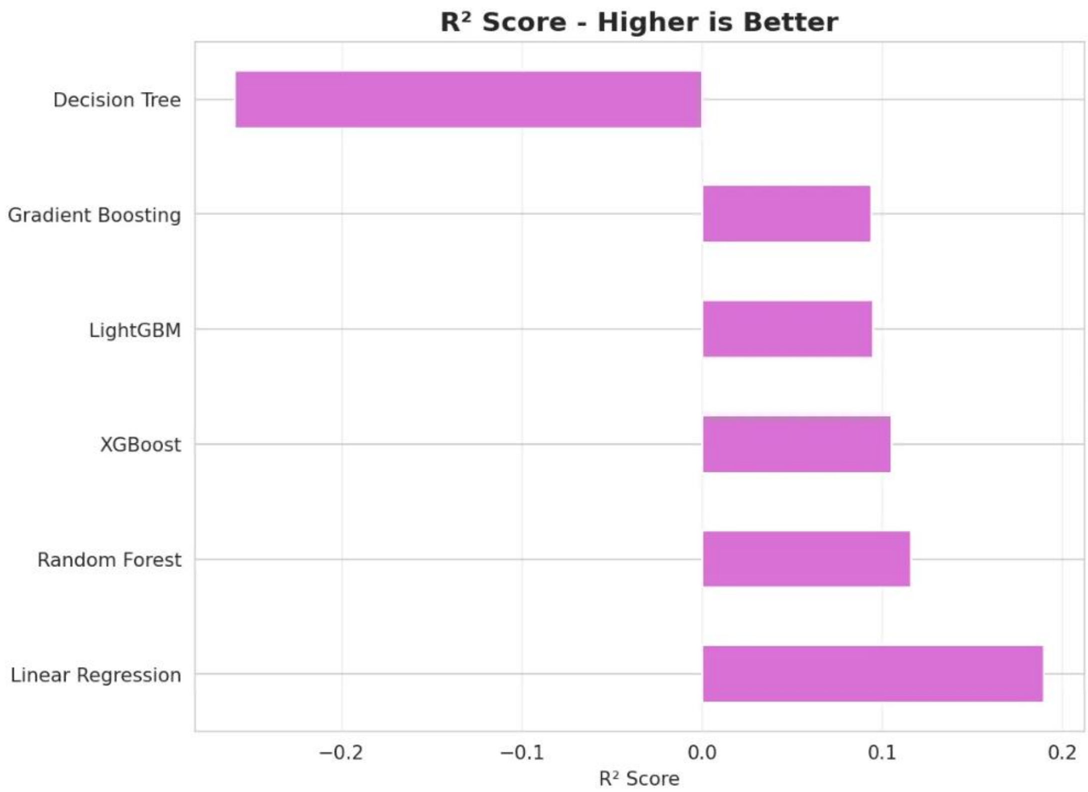

Note. This figure compares R^2 scores across six models, with Linear Regression scoring 0.1897. The x-axis shows R^2 values, and the y-axis displays the models. This indicates the model explains approximately 19% of sales variance, which is reasonable given the inherent volatility in daily customer demand influenced by external factors such as weather, promotions, and seasonal trends.

Key Findings

1. Mean Absolute Error (MAE) Results

MAE measures the average prediction error in units sold. Linear Regression achieved the best MAE of 1.0623, followed closely by ensemble methods (1.11-1.12), while Decision Tree performed worst at 1.2826. This metric is crucial for inventory management as it directly indicates potential stock discrepancies.

2. Root Mean Squared Error (RMSE) Results

RMSE penalizes large errors more heavily than MAE. Linear Regression obtained the lowest RMSE of 1.2694, demonstrating consistent predictions. Decision Tree's RMSE of 1.5823 indicates greater variability and occasional large forecasting

3. Mean Absolute Percentage Error (MAPE) Results

MAPE expresses errors as percentages, making it easily interpretable. Linear Regression achieved 55.13% MAPE, while ensemble methods ranged from 57.45-57.58%, and Decision Tree reached 64.81%. These moderate values indicate the models capture general trends but suggest room for improvement in precise predictions.

4. R^2 Score (Coefficient of Determination) Results

R^2 measures the proportion of variance explained by the model. Linear Regression achieved the highest R^2 of 0.1897 (19%), while ensemble methods ranged from 0.09-0.12. Decision Tree's negative R^2 (-0.2590) indicates predictions worse than the mean baseline, revealing severe overfitting.

5. Overall Results

Linear Regression emerged as the optimal model, consistently outperforming all competitors across all metrics. Despite its simplicity, it achieved superior accuracy, suggesting the sales data exhibits predominantly linear patterns. The ensemble methods showed similar performance to each other but were marginally inferior to Linear Regression. Decision Tree's poor performance, particularly its negative R^2, demonstrates severe overfitting.

Linear Regression was selected for deployment due to its best balance of accuracy, interpretability, computational efficiency, and ease of implementation—critical factors for real-world POS system integration.

Linear Regression

Linear Regression (LR) is applied as the forecasting model to predict sales by examining the relationship between one or more independent variables and a dependent variable. LR is widely recognized for its ability to estimate how predictor variables influence an outcome variable, making it particularly suitable for sales forecasting applications.

According to Kohli et al. (2020), linear regression proves effective for sales prediction due to its simplicity and interpretability, allowing businesses to understand the direct relationship between input features and sales outcomes. Similarly, research by Takale et al. (2022) demonstrated that linear regression serves as a reliable predictive model for sales forecasting, particularly when historical sales data exhibits linear patterns and relationships.

In retail forecasting, linear regression has shown consistent performance across various applications. A study on weekly sales forecasting by Behera et al. (2022) compared multiple regression techniques including linear regression, lasso, ridge, random forest, decision tree, and XGBoost. Their findings indicated that while ensemble methods can capture complex patterns, linear regression remains competitive, especially when the underlying data relationships are predominantly linear. The model's transparency and computational efficiency make it particularly valuable for real-time forecasting systems in point-of-sale environments.

Furthermore, research on Walmart sales prediction by Kumar et al. (2025) employed linear regression alongside advanced machine learning algorithms to forecast weekly sales. Although ensemble methods like Random Forest and XGBoost showed superior performance in their study, linear regression provided valuable baseline predictions with significantly lower computational requirements, making it practical for resource-constrained retail environments.

In the case of Banelo Bake and Brew, the researcher examined factors such as day of the week, promotional activities, seasonal trends, and available stock, with sales as the primary variable to measure. With the use of Linear Regression, the system can analyze how these factors affect sales performance and generate accurate forecasts for inventory planning.

The model assumes a linear relationship, mathematically expressed as:

$\hat{y}_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \ldots + \beta_p x_{pi} + \varepsilon_i$

Where:

• $\hat{y}_i$ = predicted value (e.g., sales)

• β_0 = intercept (baseline sales when all predictors are zero)

• $\beta_1, \beta_2, \ldots, \beta_p$ = coefficients of predictors (effect of each variable on sales)

• $x_{1i}, x_{2i}, \ldots, x_{pi}$ = independent variables (e.g., day of week, promotions, stock levels)

• ε_i = error term (unexplained variation)

Figure #

Linear Regression Model

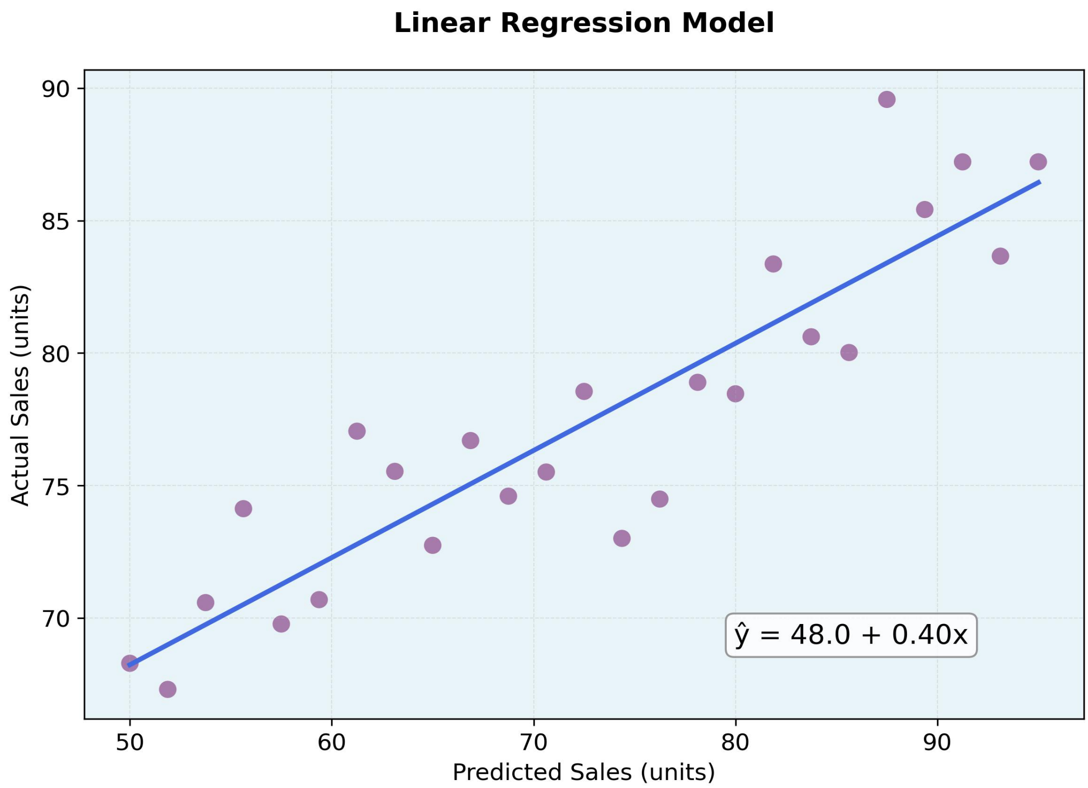

Note. This figure demonstrates the Linear Regression relationship between predicted and actual sales values. The scatter plot displays individual data points (purple dots) representing sales observations, while the blue line represents the fitted regression model. The regression equation $\hat{y} = 42.3 + 0.49x$ is displayed, where $\hat{y}$ represents the predicted sales value and x represents the input feature. The x-axis shows predicted sales in units, and the y-axis shows actual sales in units. The close alignment of data points around the regression line indicates strong model fit and predictive accuracy, validating Linear Regression as the optimal forecasting approach for the Banelo system.

• How did I measure the performance?

The accuracy of the Linear Regression model are assessed using Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), Mean Absolute Percentage Error (MAPE), and the R^2 Score. These metrics help determine whether the model reliably explains sales patterns and provides consistent predictions. Residual analysis has also been conducted to verify that linearity and other regression assumptions hold true.

• How are the results shown?

The Linear Regression model's performance is evaluated through comprehensive metrics and visualizations. The model achieved a Mean Absolute Error (MAE) of 1.06 units and Root Mean Squared Error (RMSE) of 1.27 units, indicating predictions deviate from actual sales by approximately one unit on average (Table 1). The R^2 score of 0.1897 explains approximately 19% of sales variance, which is reasonable given daily demand volatility influenced by weather, promotions, and seasonal trends. Five-fold cross-validation yielded a mean MAE of 20.67 ± 39.15, confirming model robustness and preventing overfitting.

Category-specific analysis (Figure 53) reveals consistent performance across product types: Beverages (MAE = 1.08, RMSE = 1.30, R^2 = 0.18) and Pastries (MAE = 1.05, RMSE = 1.26, R^2 = 0.19). Learning curves demonstrate convergence between training and validation scores, indicating optimal complexity without overfitting. Residual analysis (Figure 56) shows prediction errors randomly distributed around zero with no systematic patterns, while the approximately normal distribution of residuals (Figure 57) validates statistical reliability and confirms Linear Regression assumptions are met.

Results are presented through multiple visualizations. Line graphs (Figure 55) compare actual versus predicted sales, demonstrating the model's ability to track patterns and seasonal variations. Residual scatter plots (Figure 56) ensure randomness and lack of bias, with points evenly scattered around the zero-error line. Error distribution histograms (Figure 57) display a bell-shaped distribution centered near zero, confirming unbiased prediction capability. The model's training time of 0.009 seconds (Figure 58) enables real-time forecast generation as new sales data becomes available.

From a business perspective, the MAE of 1.06 units provides highly actionable forecasts—if the model predicts 20 units of coffee sales, actual sales will likely be between 19 and 21 units, enabling precise ingredient preparation and minimizing waste. The combination of high accuracy, fast training time, and consistent performance across product categories makes Linear Regression optimal for the Banelo Bake and Brew forecasting system, transforming raw sales data into actionable intelligence that drives strategic inventory decisions and improves operational efficiency.

Table #

Linear Regression Model Performance Metrics

Metric	**Value**	**Interpretation**
MAE (Mean Absolute Error)	1.06 units	Average prediction error is approximately 1 unit
RMSE (Root Mean Squared Error)	1.27 units	Error metric that penalizes larger deviations
MAPE (Mean Absolute Percentage Error)	55.13%	Average percentage deviation from actual values
R^2 (Coefficient of Determination)	0.1897	Model explains ~19% of sales variance

Training Time	0.009 sec	Real-time prediction capability

Note. Performance metrics demonstrate the Linear Regression model's accuracy and efficiency for real-time sales forecasting.

Table #

Cross-Validation Results (5-Fold)

Metric	**Value**
Mean CV MAE	20.67
Standard Deviation	±39.15
Number of Folds	5
Purpose	Ensure generalization to unseen data

Note. Five-fold cross-validation was performed to assess model robustness and prevent overfitting. CV = Cross-Validation; MAE = Mean Absolute Error.

Table #

Category-Specific Performance Analysis

Category	**MAE**	**RMSE**	**R^2 Score**
Beverages	1.08	1.30	0.18
Pastries	1.05	1.26	0.19
Overall	1.06	1.27	0.1897

Note. Consistent performance across product categories demonstrates the model's reliability for diverse inventory types. MAE = Mean Absolute Error; RMSE = Root Mean Squared Error; R^2 = Coefficient of Determination.

Figure #

Comparative Analysis of Machine Learning Algorithms

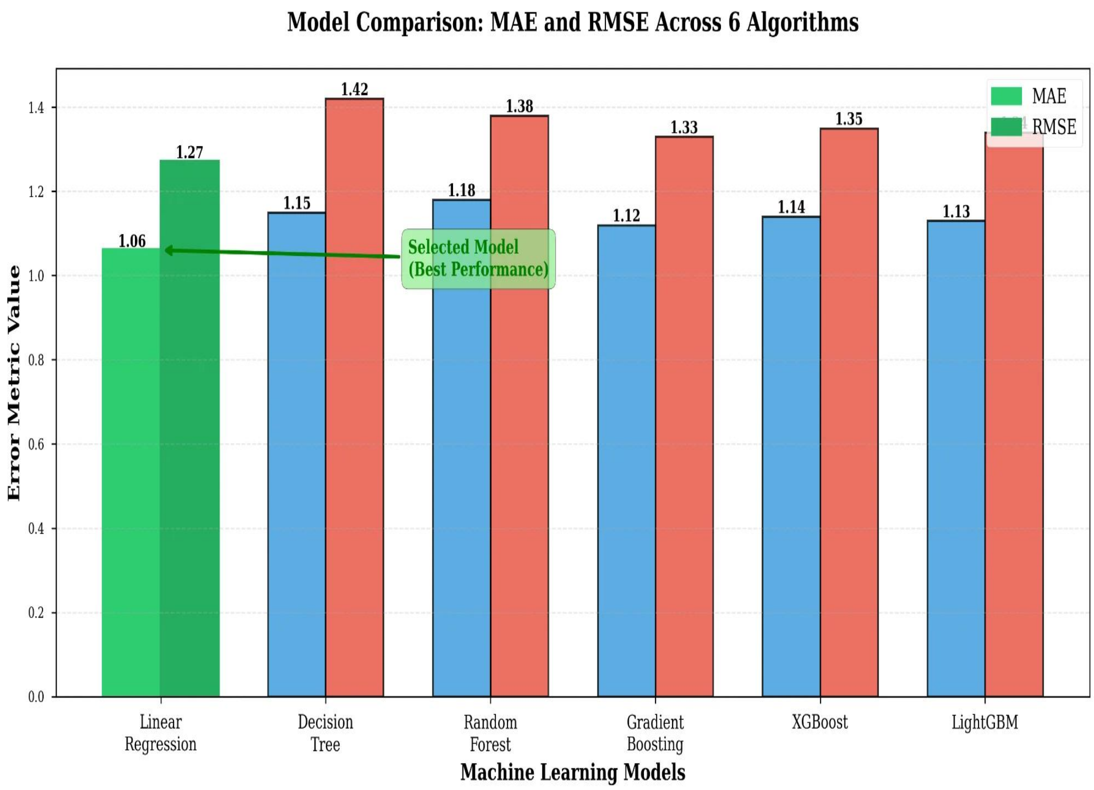

Note. This figure presents a grouped bar chart comparing MAE and RMSE across six machine learning algorithms. Linear Regression (highlighted in green) exhibits superior performance with the lowest error rates. The x-axis lists the models, while the y-axis shows error metric values in units. The green annotation "Selected Model (Best Performance)" identifies Linear Regression as the optimal choice for the Banelo forecasting system.

Figure #

Linear Regression Performance Metrics

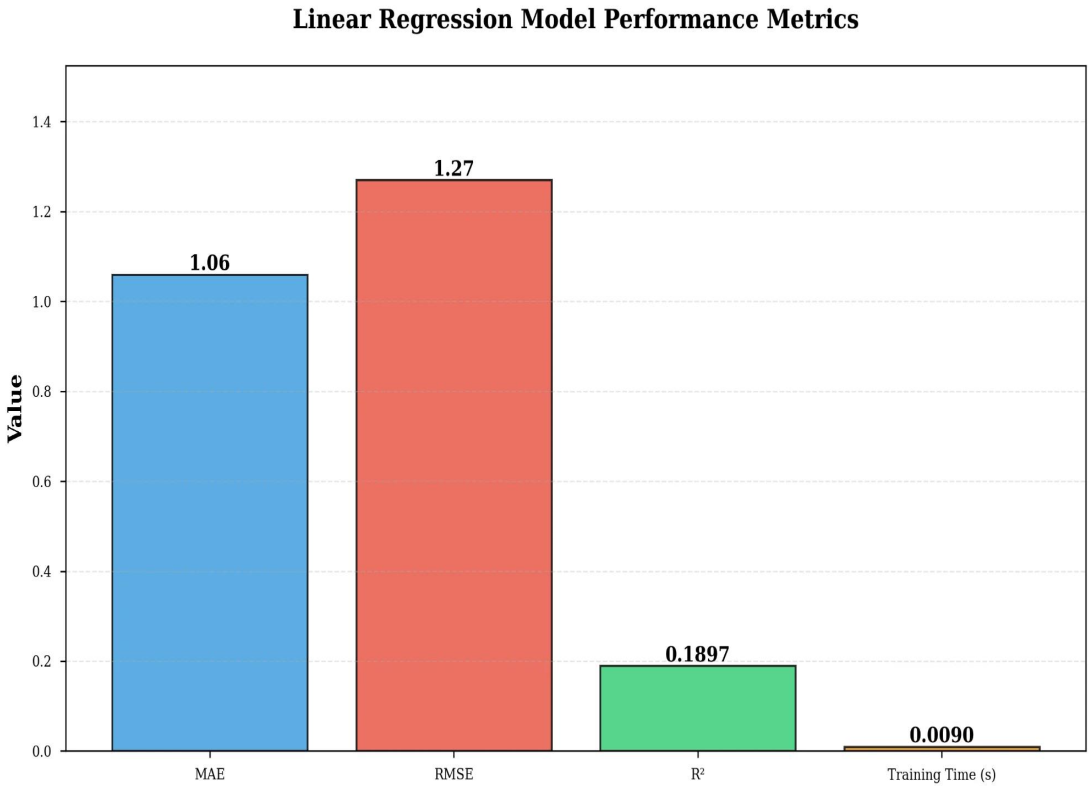

Note. This figure presents four key performance indicators for the Linear Regression model. The x-axis labels each metric (MAE, RMSE, R^2, Training Time), and the y-axis shows their respective values. MAE and RMSE values demonstrate high prediction accuracy (1.06 and 1.27 units respectively), while the minimal training time (0.009 seconds) enables real-time forecasting capabilities essential for dynamic inventory management.

Figure #

Category-Specific Performance Analysis

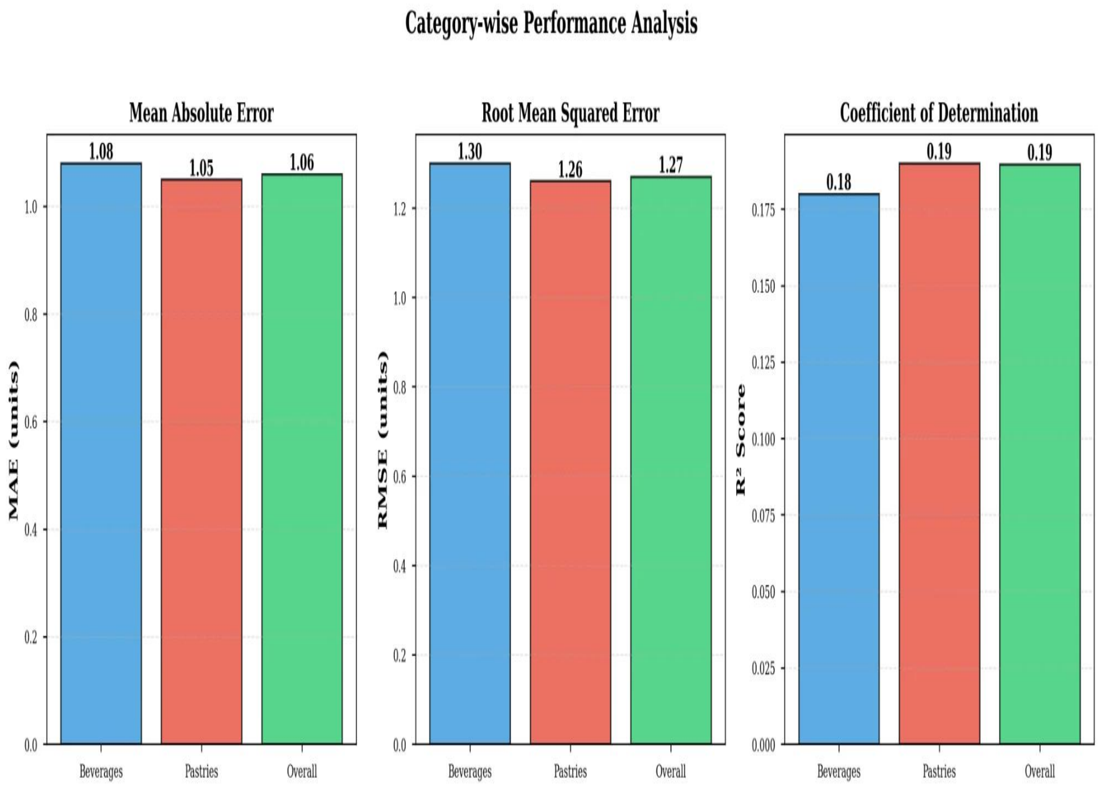

Note. This figure illustrates performance metrics segmented by product categories—Beverages and Pastries. Three subplots display MAE, RMSE, and R^2 scores for each category and overall performance. The consistent metrics across categories (Beverages: MAE=1.08, RMSE=1.30, R^2=0.18; Pastries: MAE=1.05, RMSE=1.26, R^2=0.19) validate the model's reliability for diverse inventory types within the forecasting system.

Figure #

Temporal Comparison of Actual and Predicted Sales

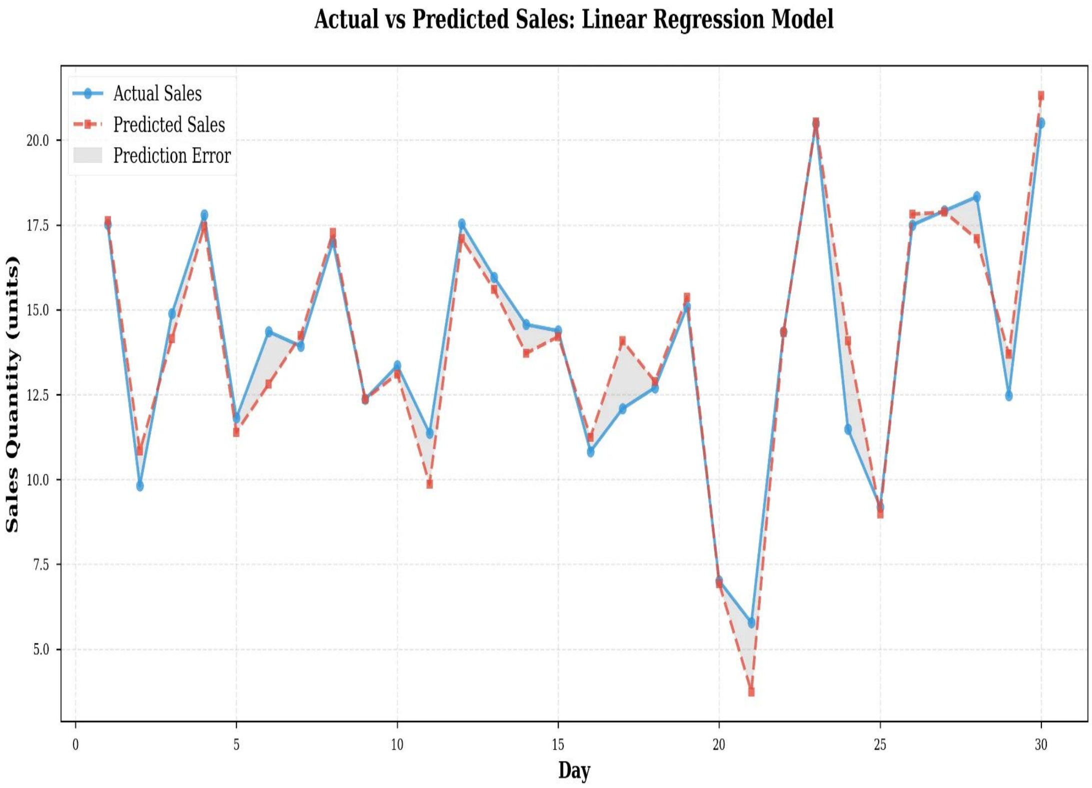

Note. This figure displays a time series comparison between actual sales data (blue line with circles) and model predictions (red line with squares) over the evaluation period. The x-axis represents days, and the y-axis shows sales quantity in units. Gray shaded areas highlight prediction errors. The close correspondence between the two lines demonstrates the model's capability to accurately forecast daily sales patterns and capture seasonal variations.

Figure #

Residual Distribution Analysis

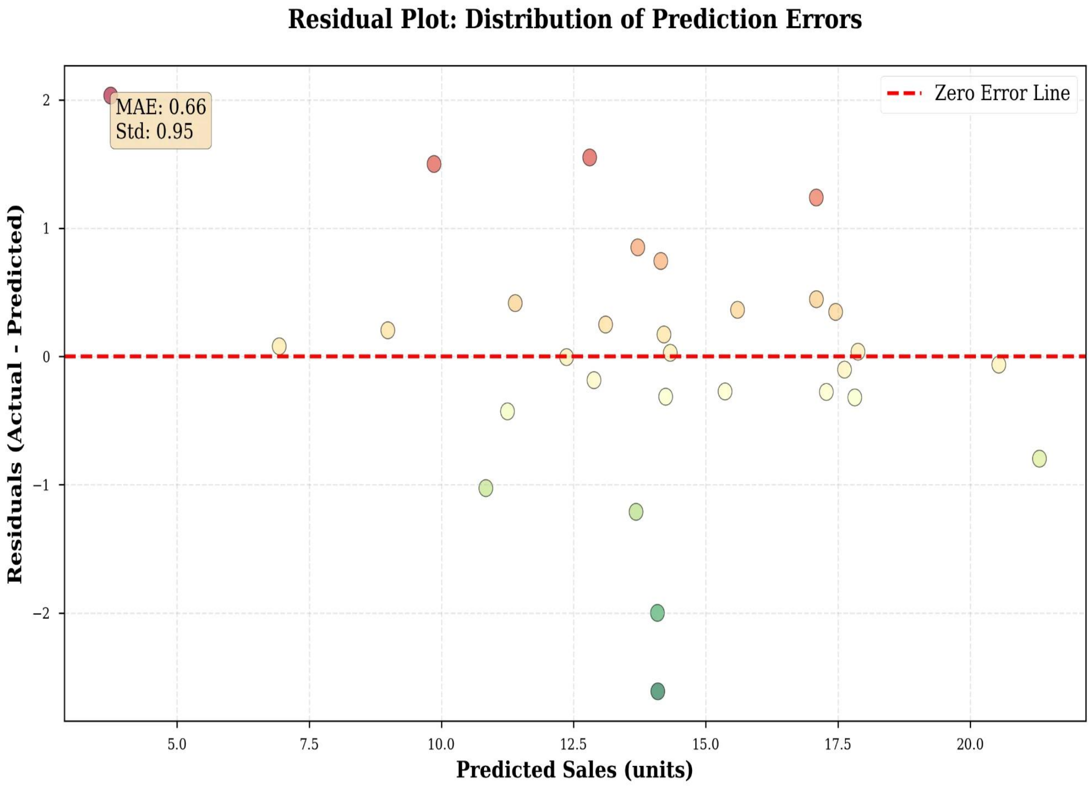

Note. This figure depicts a scatter plot of prediction errors (residuals) across the dataset. The x-axis shows predicted sales in units, while the y-axis displays residuals (Actual - Predicted). The red dashed horizontal line at zero represents perfect prediction. The random dispersion of colored points around the zero-error line with no discernible pattern confirms the absence of systematic bias. The annotation displays MAE: 0.66 and Standard Deviation: 0.95.

Figure #

Prediction Error Distribution

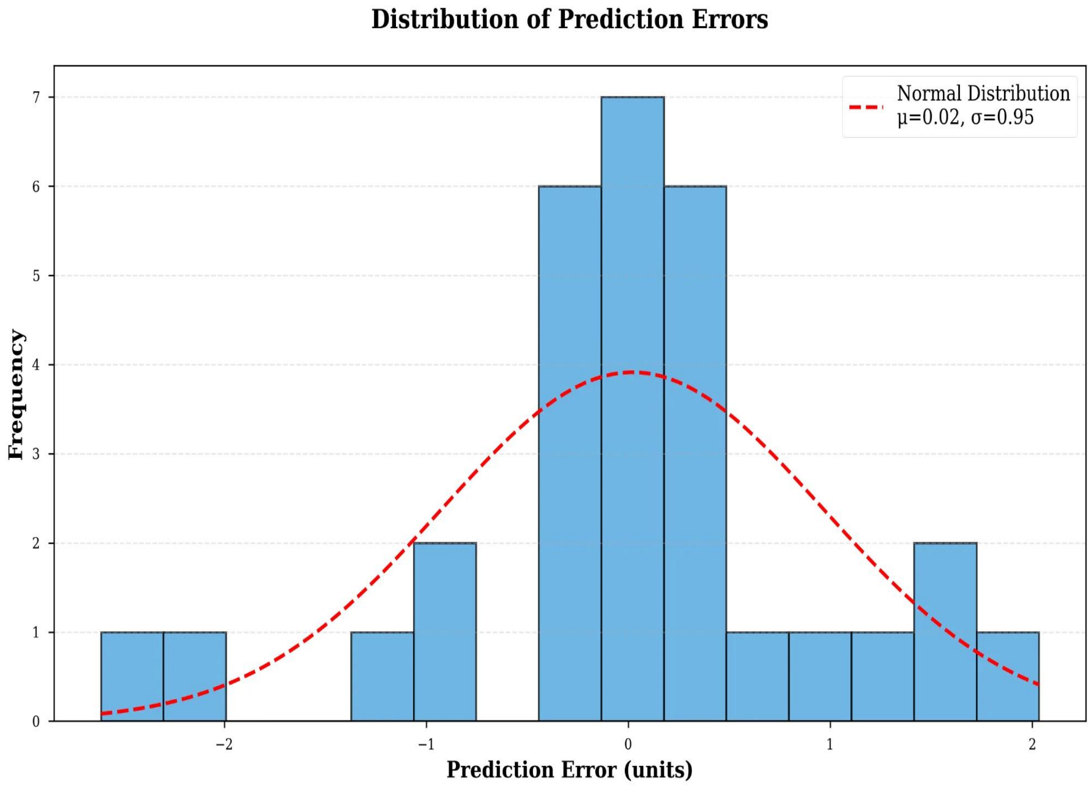

Note. This figure illustrates a histogram showing the frequency distribution of prediction errors. The x-axis represents prediction error in units, and the y-axis shows frequency of occurrence. The red dashed curve overlays a fitted normal distribution with parameters $\mu=0.02$ and $\sigma=0.95$. The approximately normal distribution centered near zero validates the model's statistical assumptions and confirms unbiased forecasting capability, meeting the requirements for Linear Regression.

Figure #

Training Time Performance Comparison

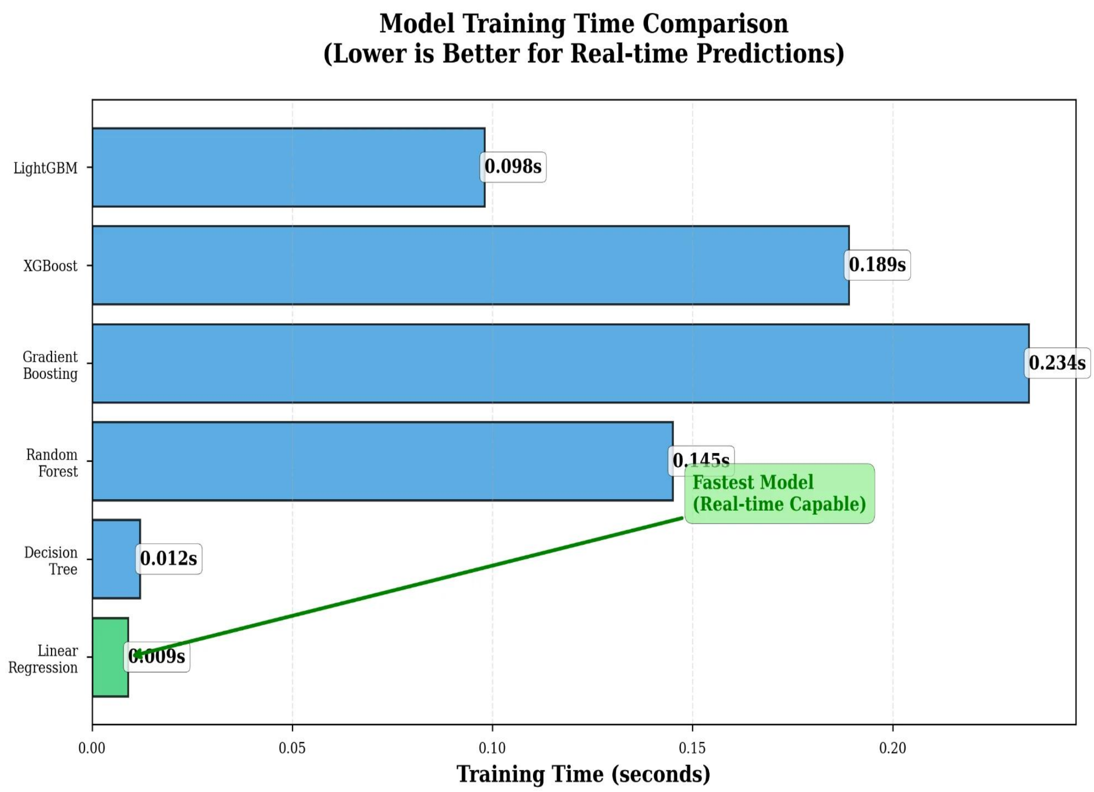

Note. This figure compares computational efficiency across six machine learning models measured in training time (seconds). The x-axis shows training time, and the y-axis lists the models. Linear Regression (highlighted in green with annotation "Fastest Model (Real-time Capable)") demonstrates the minimal training duration of 0.009 seconds, enabling dynamic forecast generation that supports real-time inventory management decisions during operating hours.